



# Convolutional Neural Networks

---

Machine Learning Decal

Hosted by Machine Learning at Berkeley



# Agenda

Motivation

Convolutions

Convolutions in Neural Networks

Questions

Demo

## Motivation

---

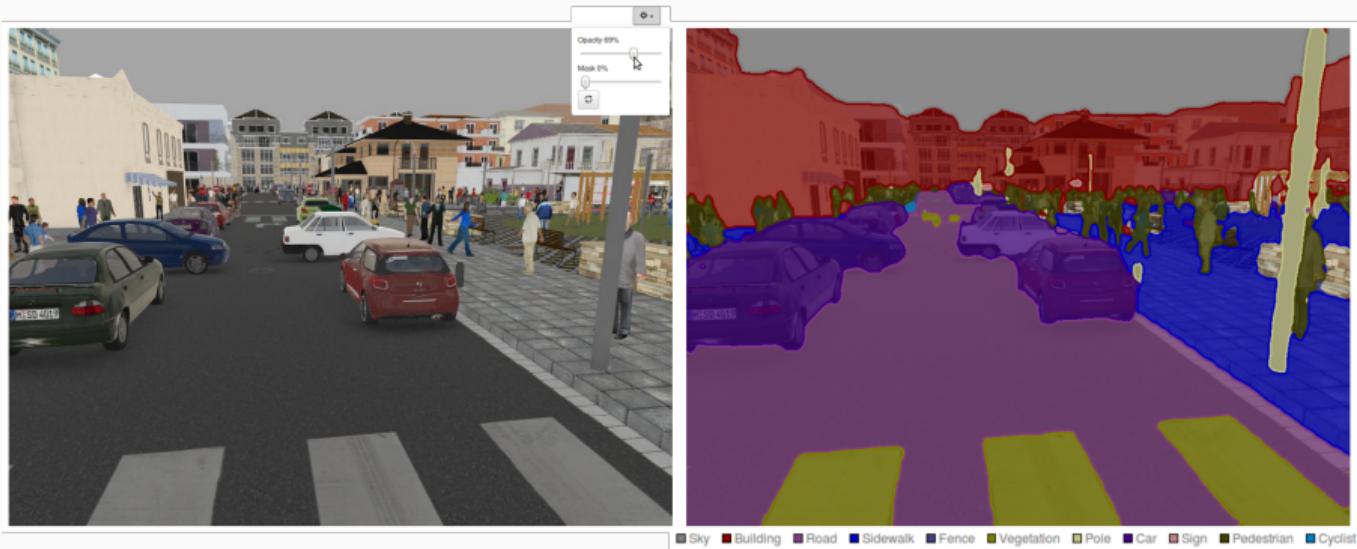
# Applications of CNNs



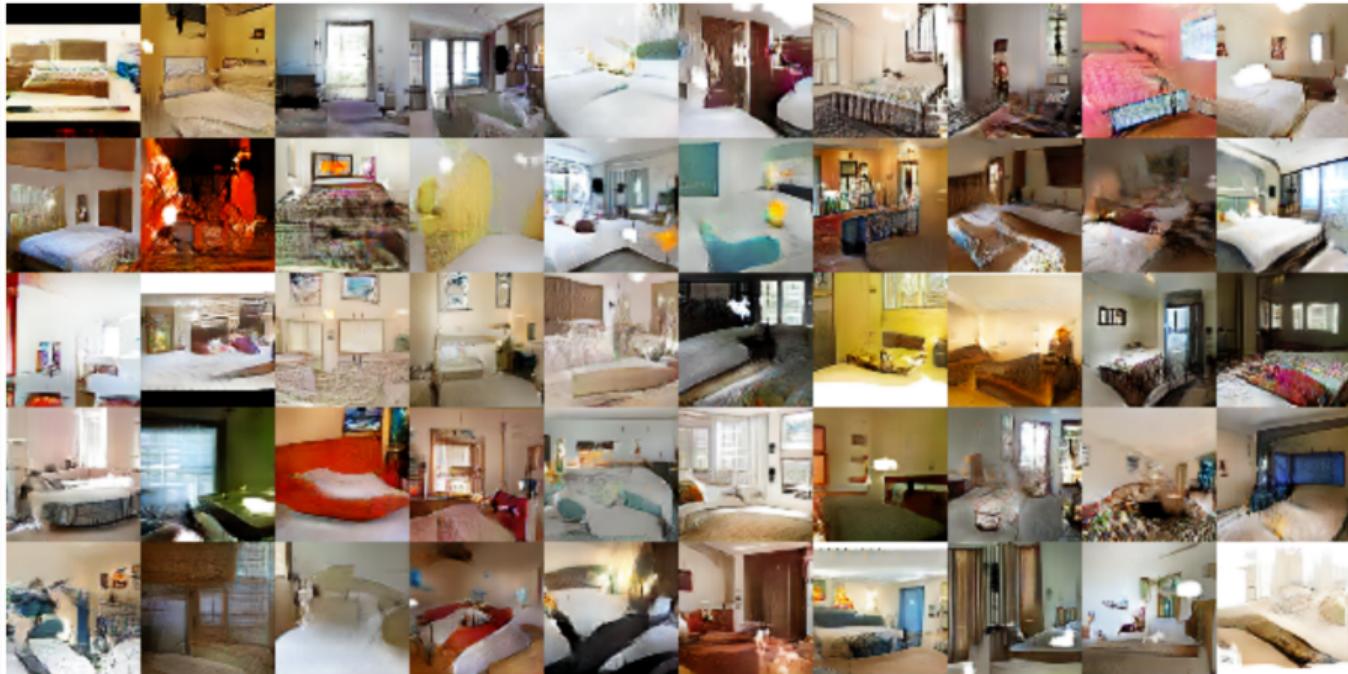
- Computer vision
  - Face recognition
  - Scene labeling
  - Image segmentation
  - Image classification
  - Action recognition
  - Human pose estimation
  - Playing Games (Atari games, Pacman)
  - Medical image analysis



# Image Segmentation



# Image Generation



# Image Style Transfer

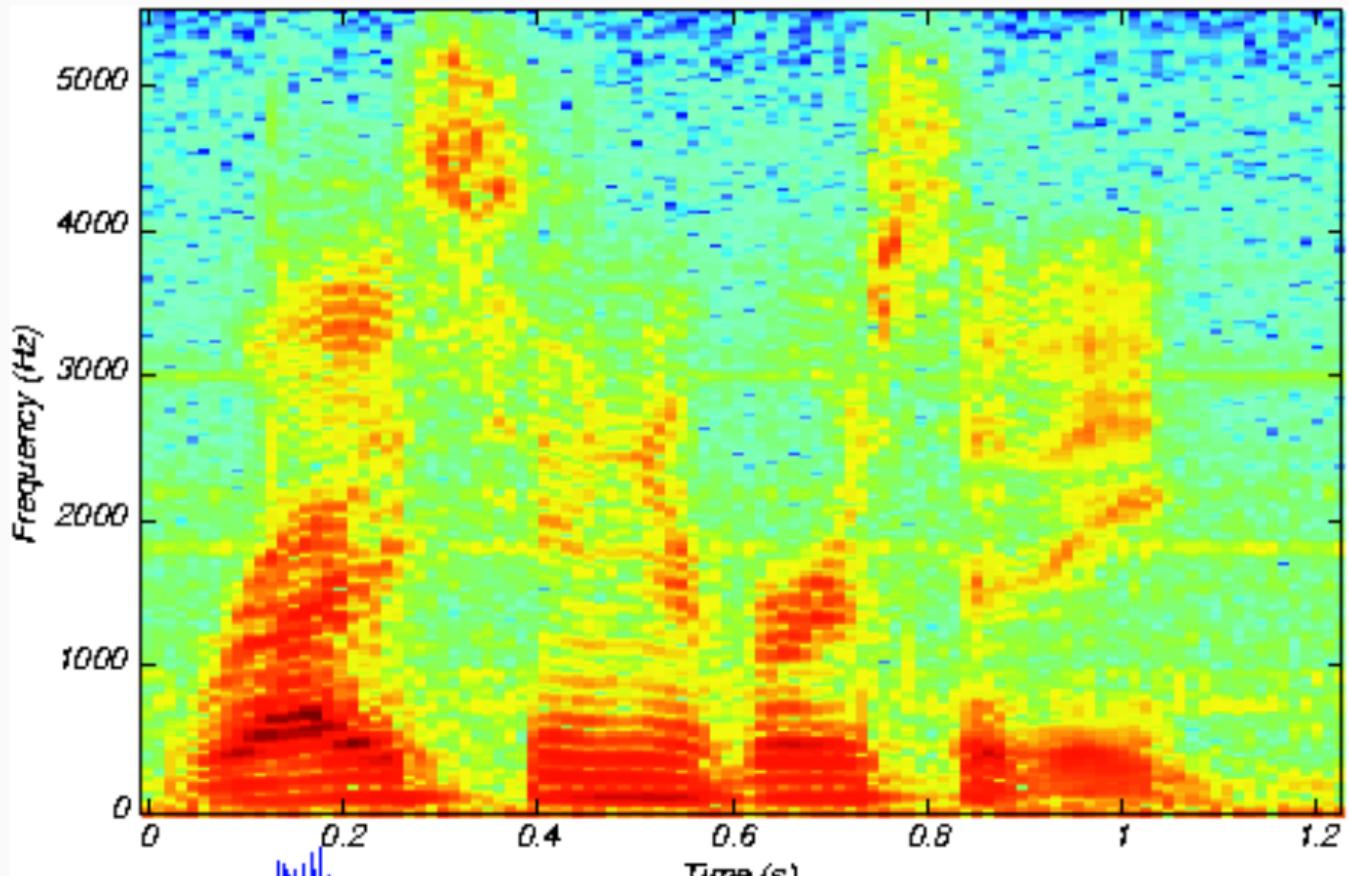


## Another Style Transfer Example



Cycle GAN: <https://github.com/junyanz/CycleGAN>

# Audio Processing

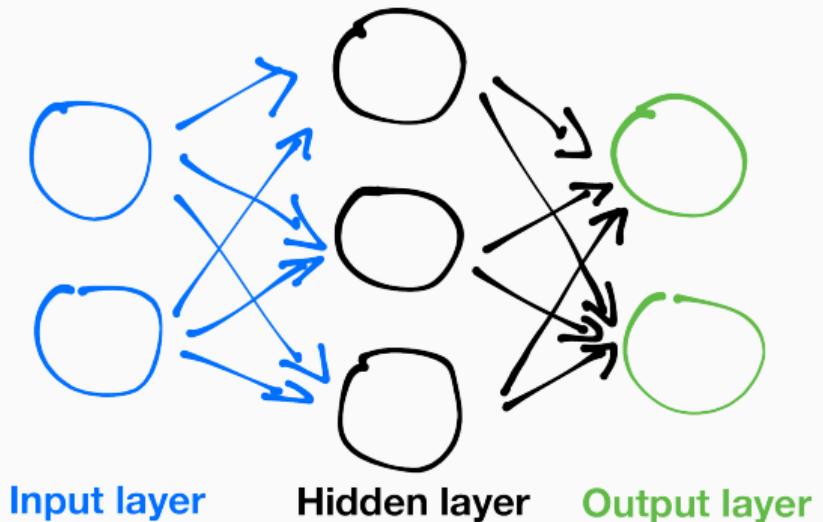


## Industry use of CNNs?



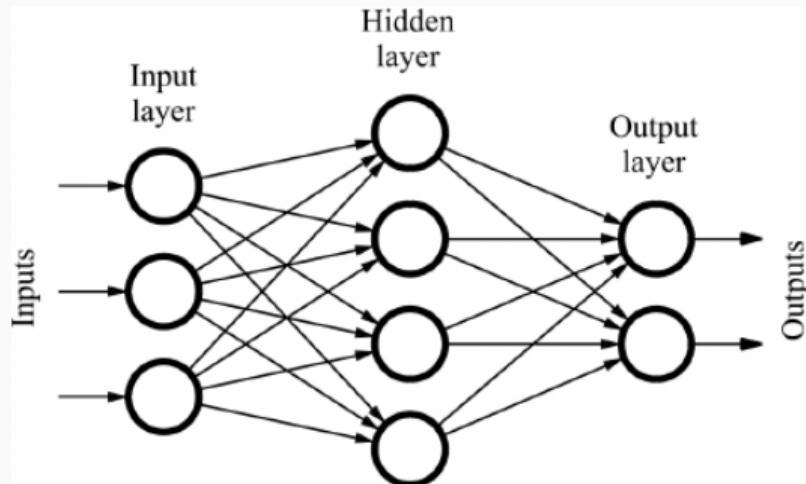
Genentech

# Review: Feed Forward Neural Network



- Layered Architecture
- Three types of layers:
  - **Input Layer** - Data is passed into these neurons
  - **Hidden Layer** - These neurons are "hidden from view"
  - **Output Layer** - These neurons output the result of the network

# What's wrong with feed-forward NNs?



The curse of dimensionality.

## Problem: high-dimensional image data



- If we want a neural network to take as input  $256 \times 256$  images, the input layer will have 65,536 neurons
- ... if we have one hidden layers with 1,000 neurons, we're already at 65 million parameters
- ... if we have two hidden layers with 15,000 neurons each, we're at over a **billion** parameters

## Solution: Convolutional Neural Networks



Basic Idea: Share parameters for different parts of input.

# Convolutions

---

## Convolution Activity



I need 9 volunteers. There will be candy involved.

## 1D Convolution



First 6 people are going to be a vector.  $2, 2, 0, -2, -2, 0$ .

Other 3 people are going to be the **kernel**:  $1, -1, 1$ .

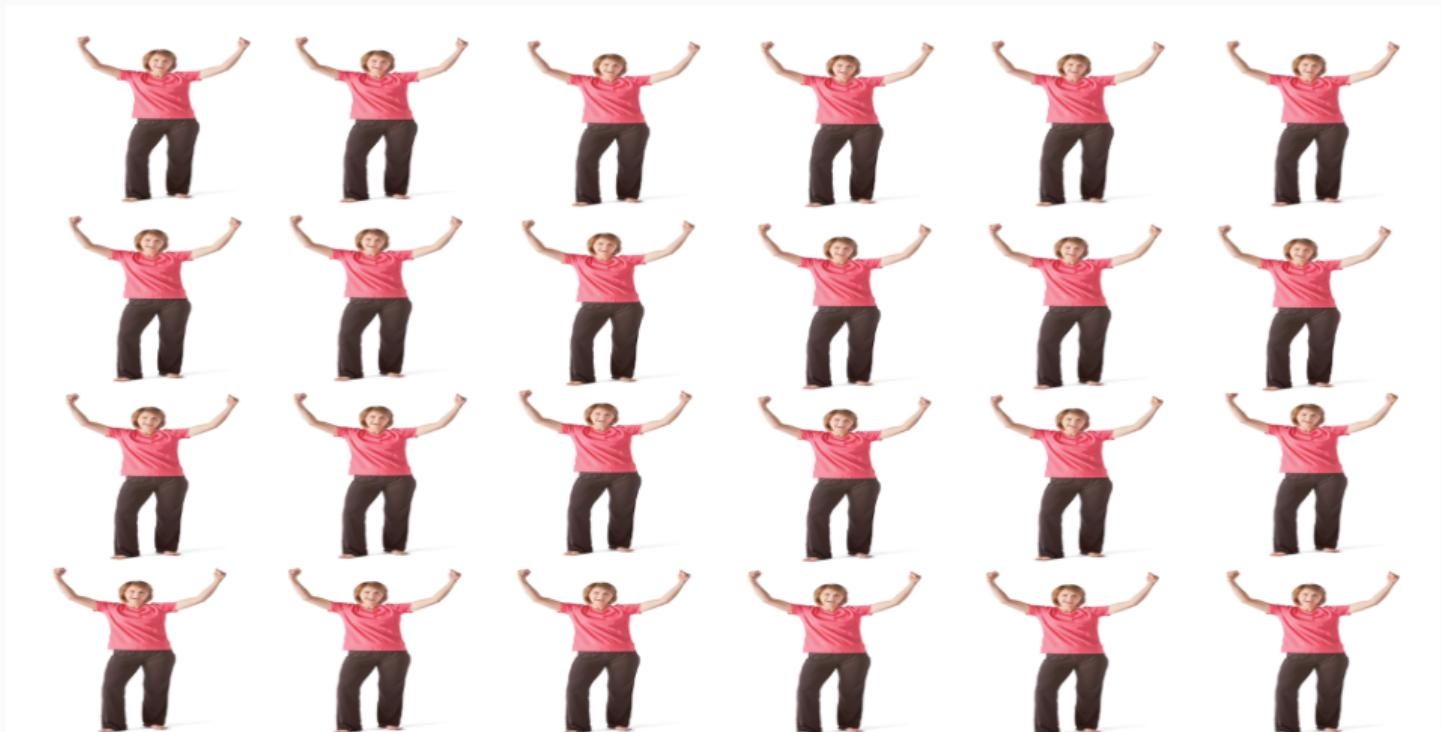
## Convolution Parameters



- Kernel Size (number of people that are moving)
- Padding Type (edge behavior)
- Stride (how much the kernel moves at each step)



## 2D Convolution



[http://deeplearning.net/software/theano/\\_images/no\\_padding\\_strides.gif](http://deeplearning.net/software/theano/_images/no_padding_strides.gif)

## 3D Convolution: Most Common in CNNs



<https://i.stack.imgur.com/FjvuN.gif>

Most Neural Network libraries call this a 2D convolution.

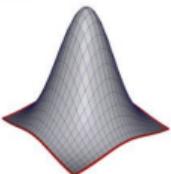
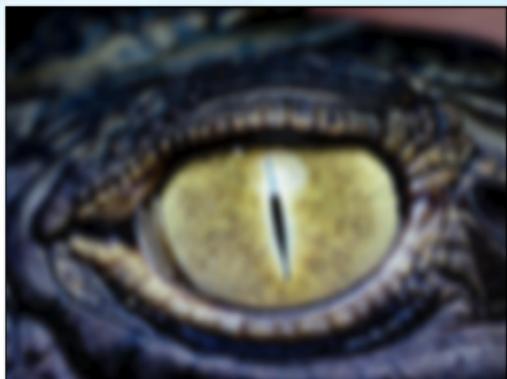
# Convolutions as filters



\*

1	2	1
2	4	2
1	2	1

=

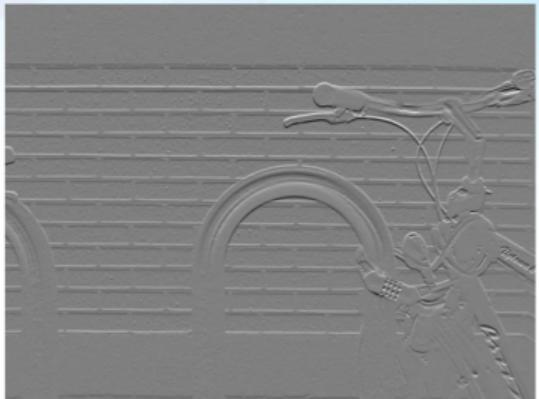




## Convolutions as filters cont.

 $*$ 

1	2	1
0	0	0
-1	-2	-1

 $=$ 

# Convolutions in Neural Networks

---

# Convolution Neural Networks

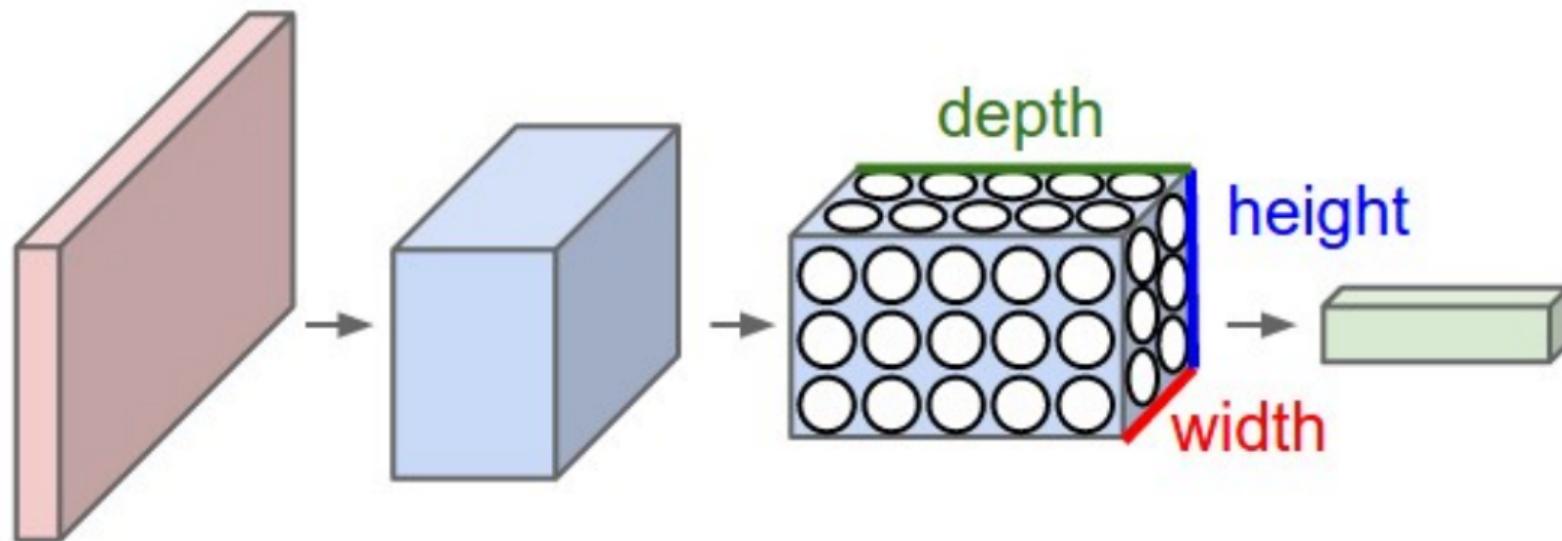


CNN = Learned convolutional filters.



## What does our input look like?

Let's say we have a  $32 \times 32$  color image. There are then 3 channels for the colors (RGB). So our input is a 3D volume of size  $32 \times 32 \times 3$



## What do our filters look like?



3D filters !!!

Our filters will then be  $n \times n \times 3$ , where  $n$  is a hyperparameter.

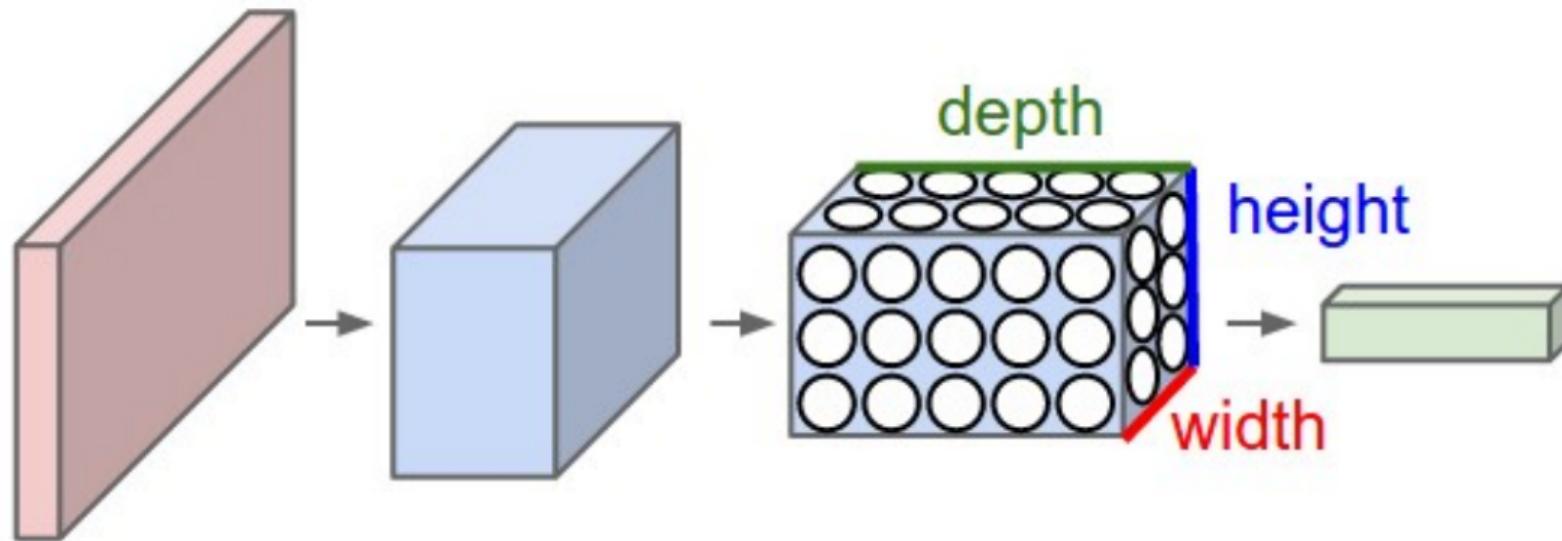
We can also have multiple filters per layer.

<https://i.stack.imgur.com/FjvuN.gif>



## What will the output of a layer look like?

Our output will also be 3D of shape  $(32 \times 32 \times m)$  where  $m$  is the number of filters in the layer.





## How do we update filters?

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$

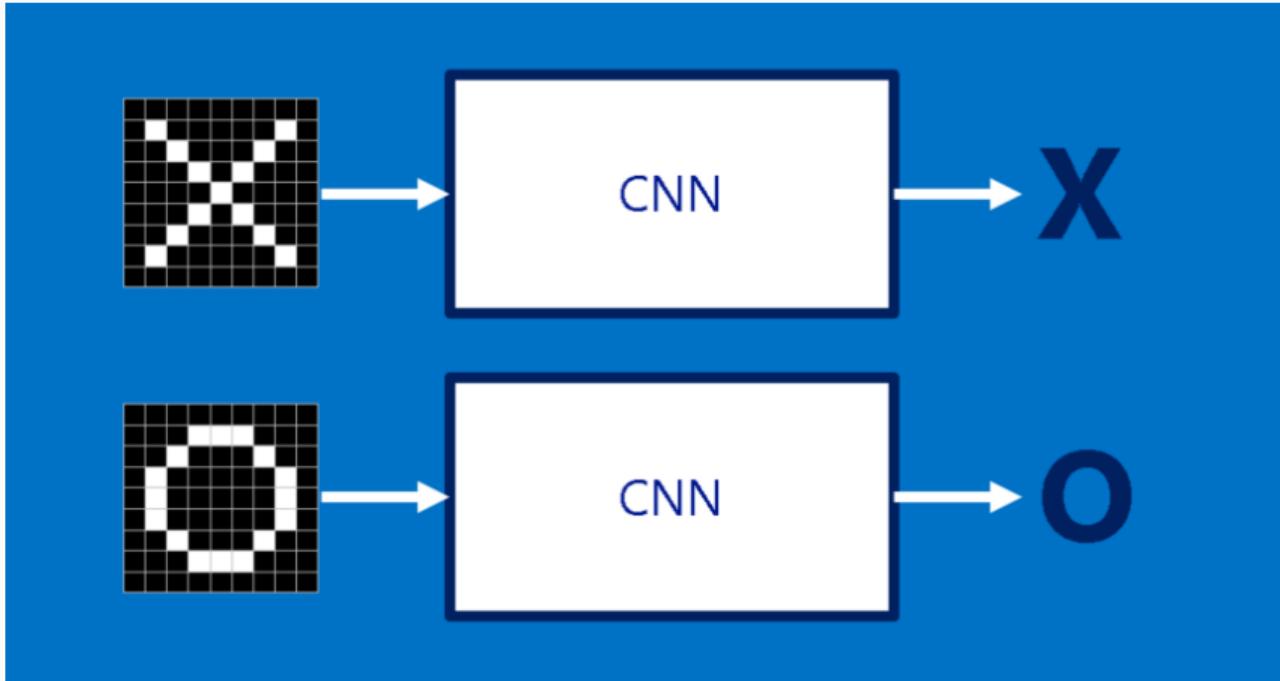
$X_{11} \partial h_{11}+$ $X_{12} \partial h_{12}+$ $X_{21} \partial h_{21}+$ $X_{22} \partial h_{22}$	$X_{12} \partial h_{11}+$ $X_{13} \partial h_{12}+$ $X_{22} \partial h_{21}+$ $X_{23} \partial h_{22}$
$X_{21} \partial h_{11}+$ $X_{22} \partial h_{12}+$ $X_{31} \partial h_{21}+$ $X_{32} \partial h_{22}$	$X_{22} \partial h_{11}+$ $X_{23} \partial h_{12}+$ $X_{32} \partial h_{21}+$ $X_{33} \partial h_{22}$

$\partial h_{11}$	$\partial h_{12}$
$\partial h_{21}$	$\partial h_{22}$

Derivative Computation (Backward pass) since pictures speak more than words



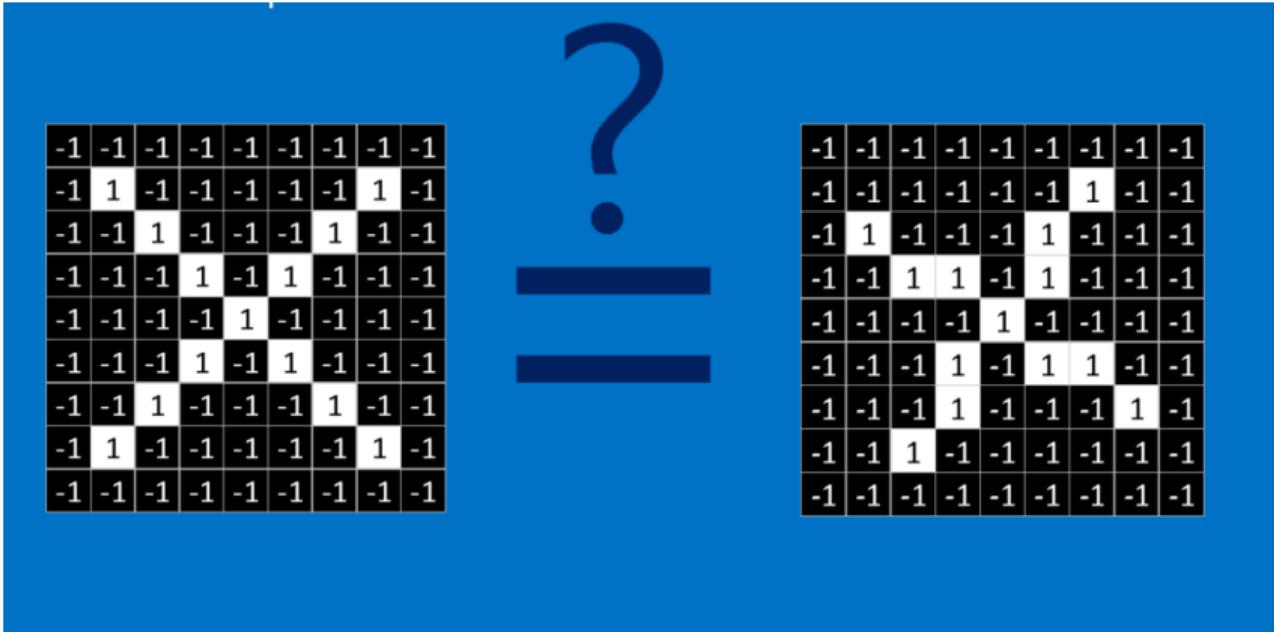
## Let's look at a toy example!



- How would a CNN classify X's vs O's?



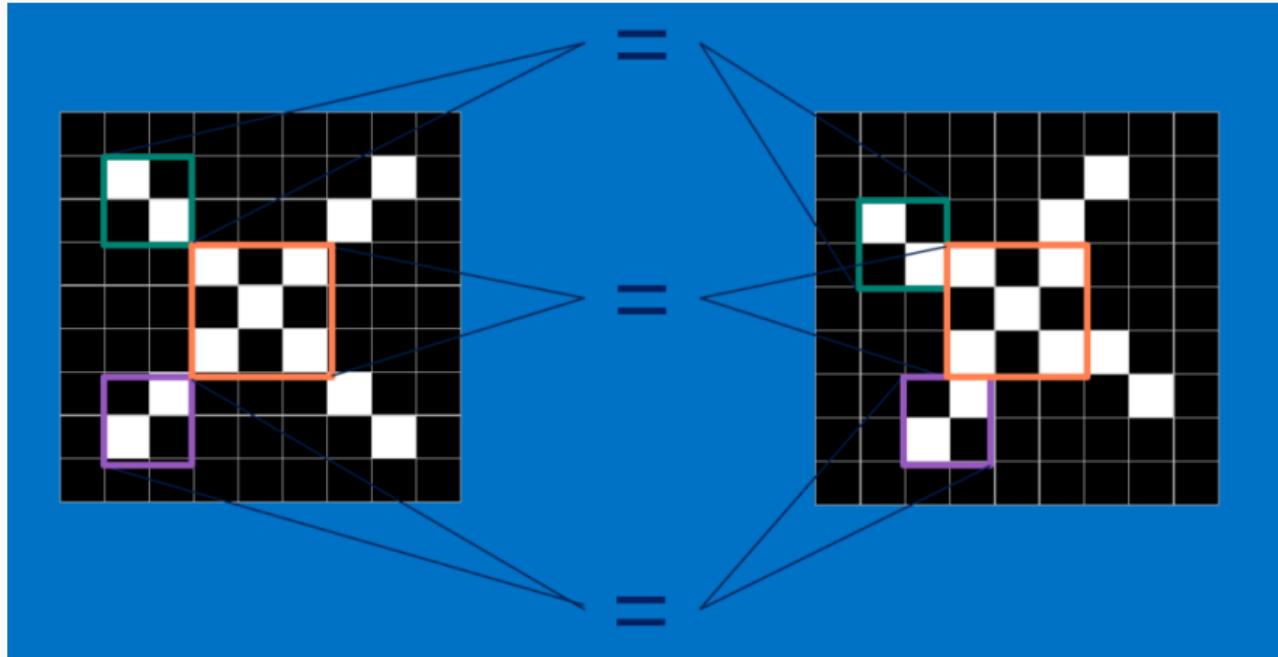
## Would you classify this as an X?



- A computer can't easily distinguish this example as an X.. so what can we do?



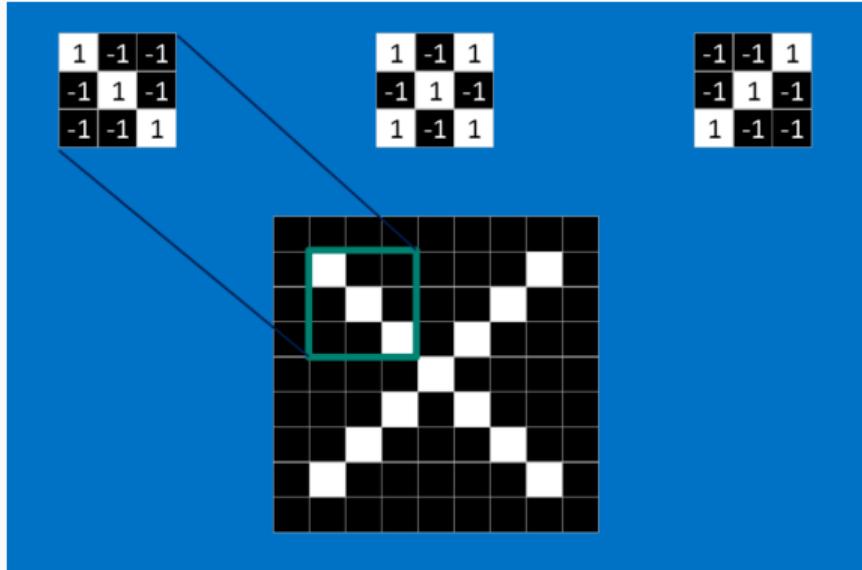
## What is a feature?



- We can break down the image into smaller images and look for patterns!
- Can we see here the similar patterns?
- These diagonals and crosses can be thought of as features!



## Convolutional Filters



- Let's break this up into smaller problems... what about smaller patterns!



## CNN Layers: Convolutions

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	1

## CNN Layers: Convolutions



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



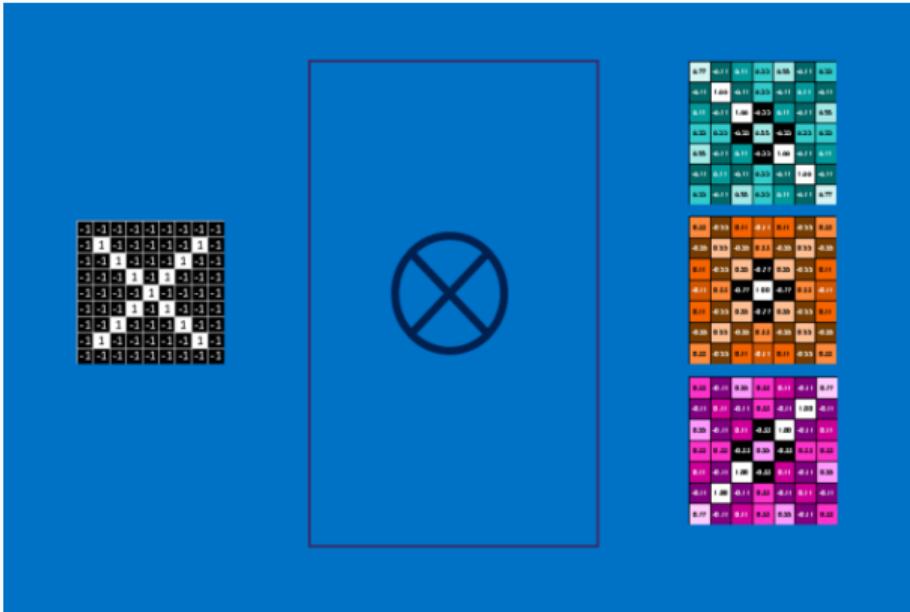
1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

- Now lets apply this process to every part of the board.
- We can look at the new board and call it a filtered version of this specific feature
- Shows us where on the image the feature matches the most

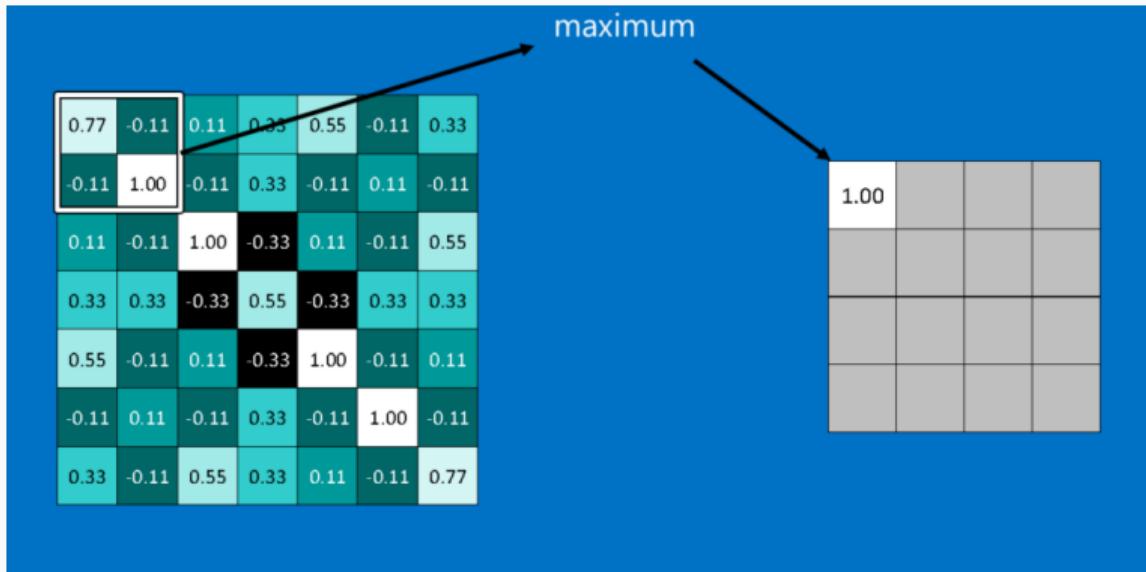
# CNN Layers: Convolutions



- Now lets apply this process to every part of the image for different features.

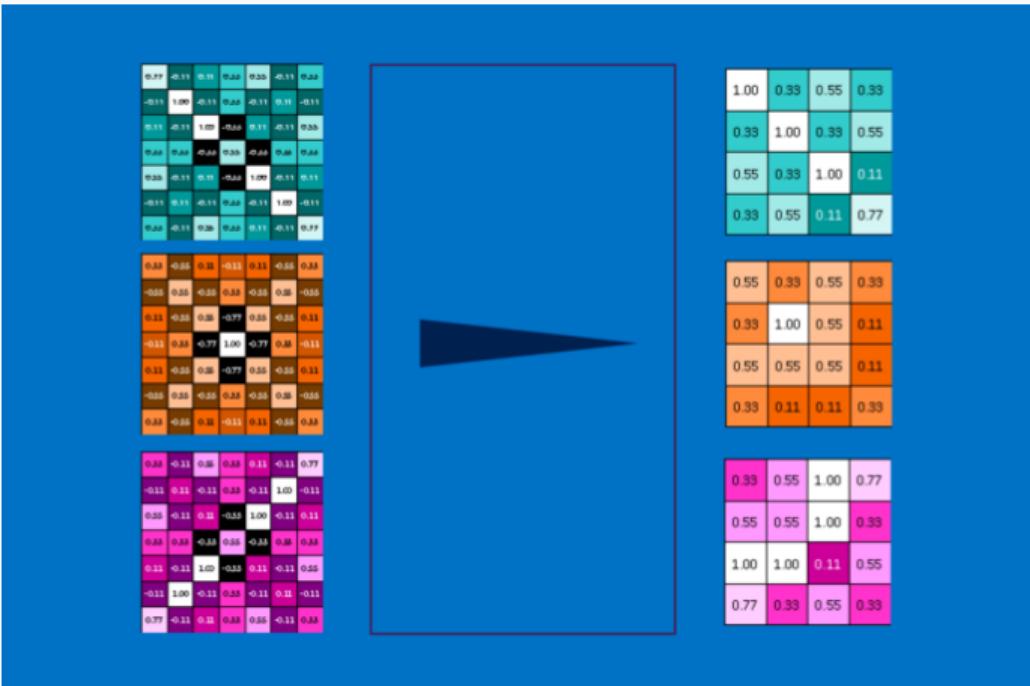


## CNN Layers: Pooling



- If CNN's are really deep, doing that math for every possible pixel gets to be quite computationally expensive
- That's why we have pooling layers, that down sample our image!

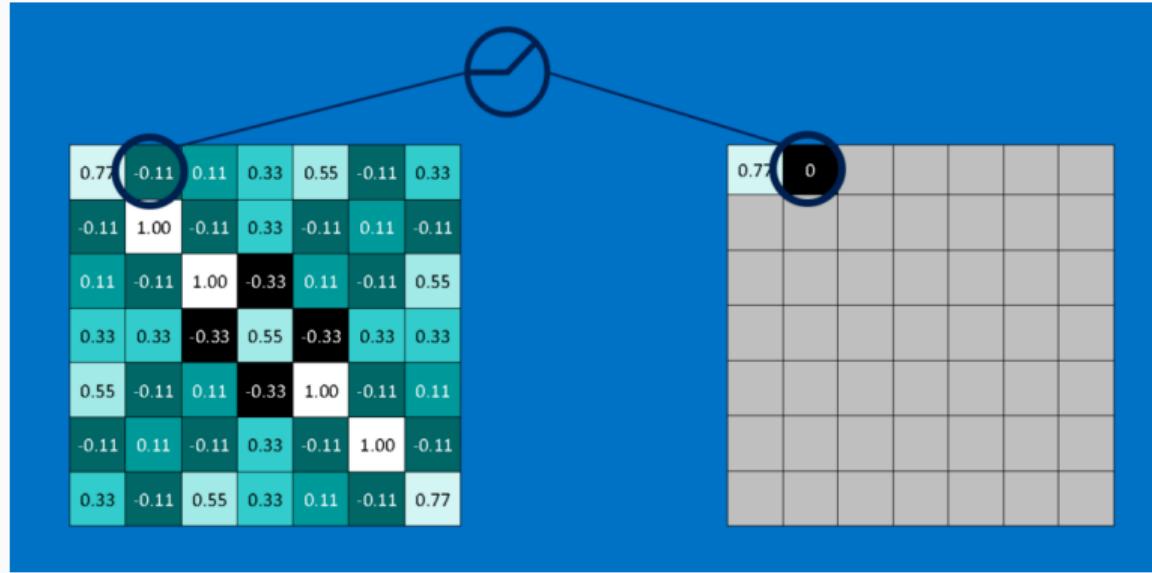
# CNN Architecture



- By applying these pooling layers, we try to capture the patterns and get rid of extra unneeded information.



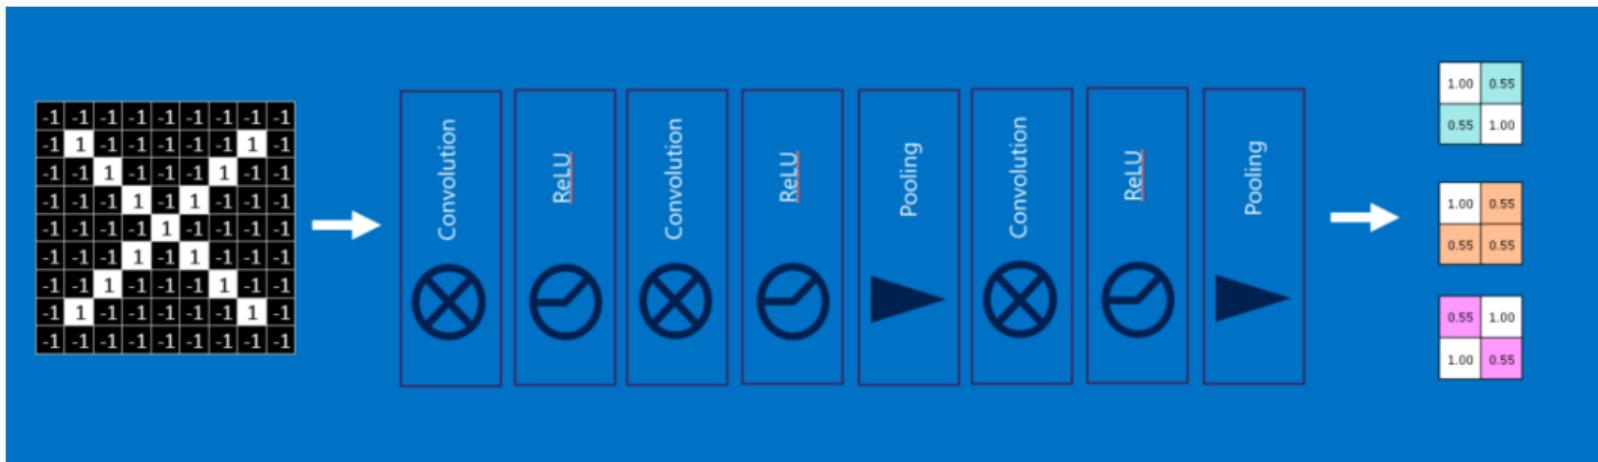
## CNN Layers: Activation



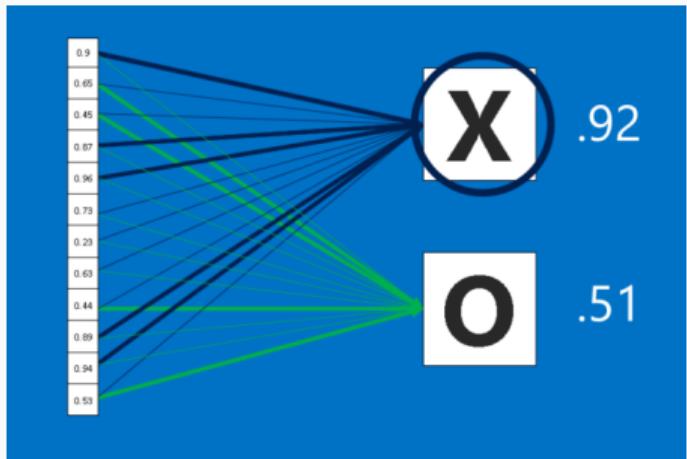
- Remember those activation functions you learned from the last lecture? We use them here too!
- Applying RELU we simply get rid of all the negatives!



# Let's put it all together!



## CNN: Fully Connected Layer

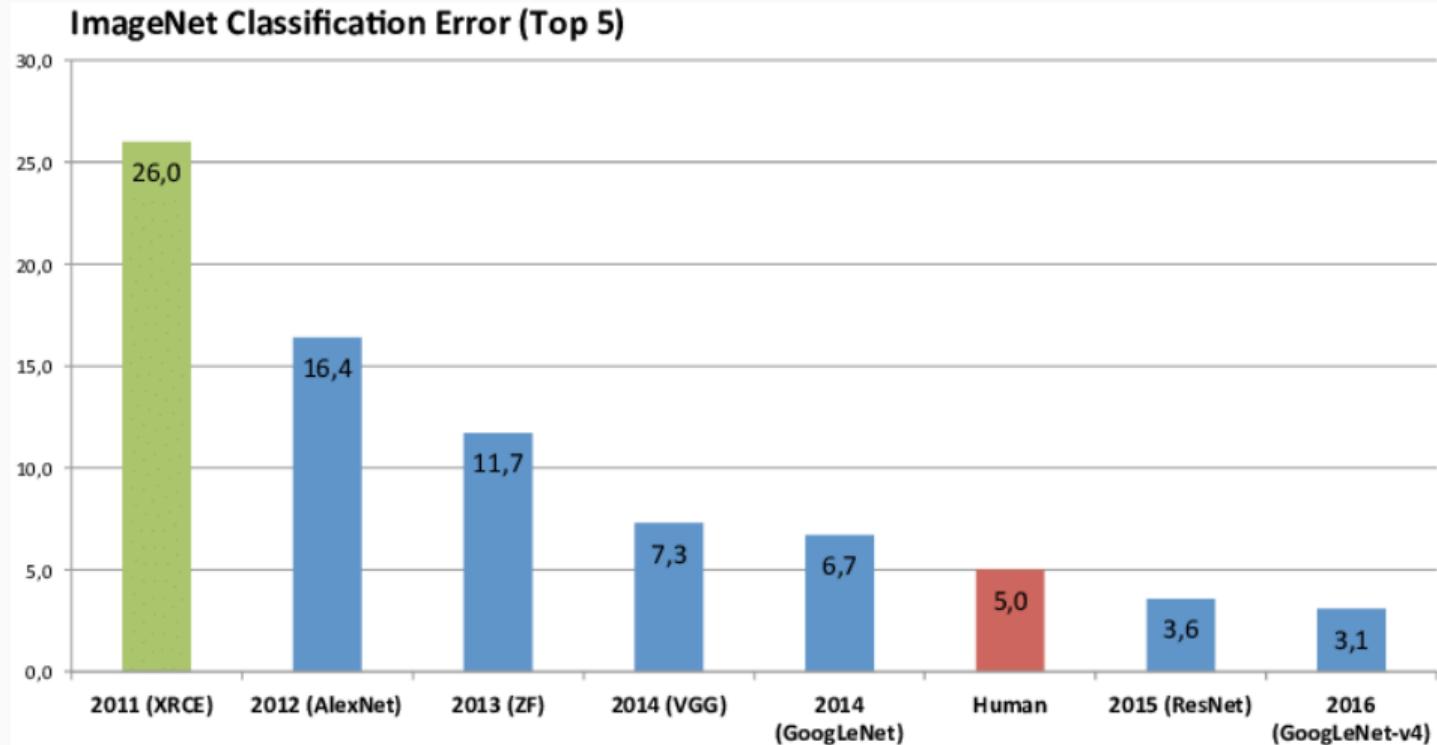


- At the end we use a fully connected layer to hold a vote if the image is either an X or O
- Each value has a weight though, so not all values contribute equally
- This goes back to the ideas you learned last week!





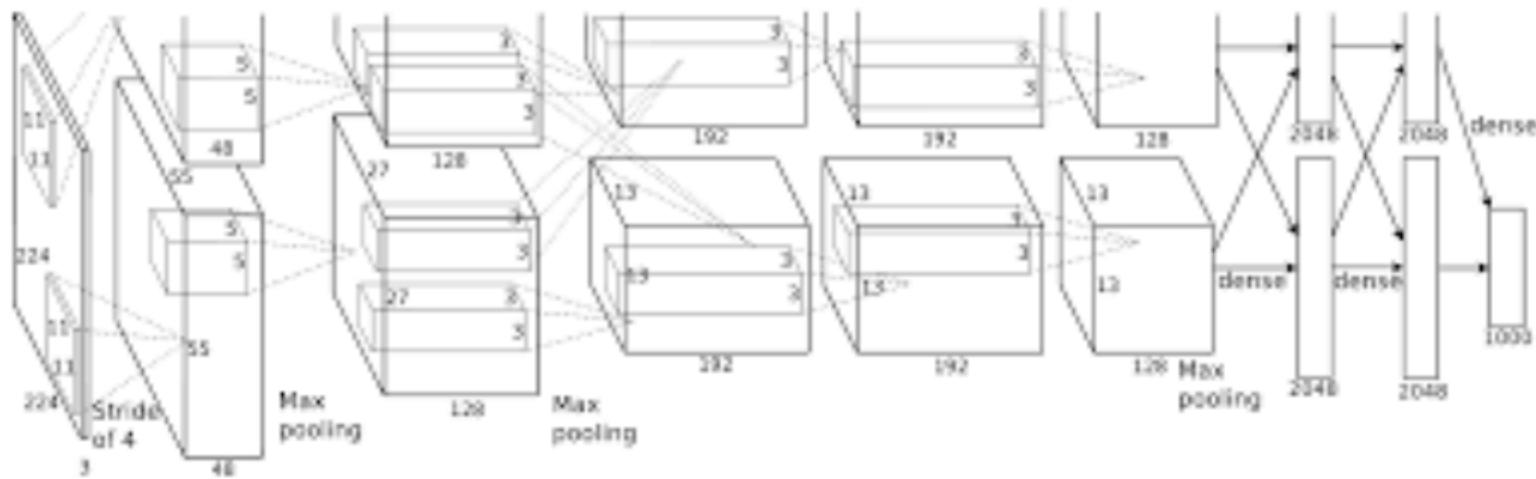
# Imagenet: Sparked Beginning of DL Image Renaissance



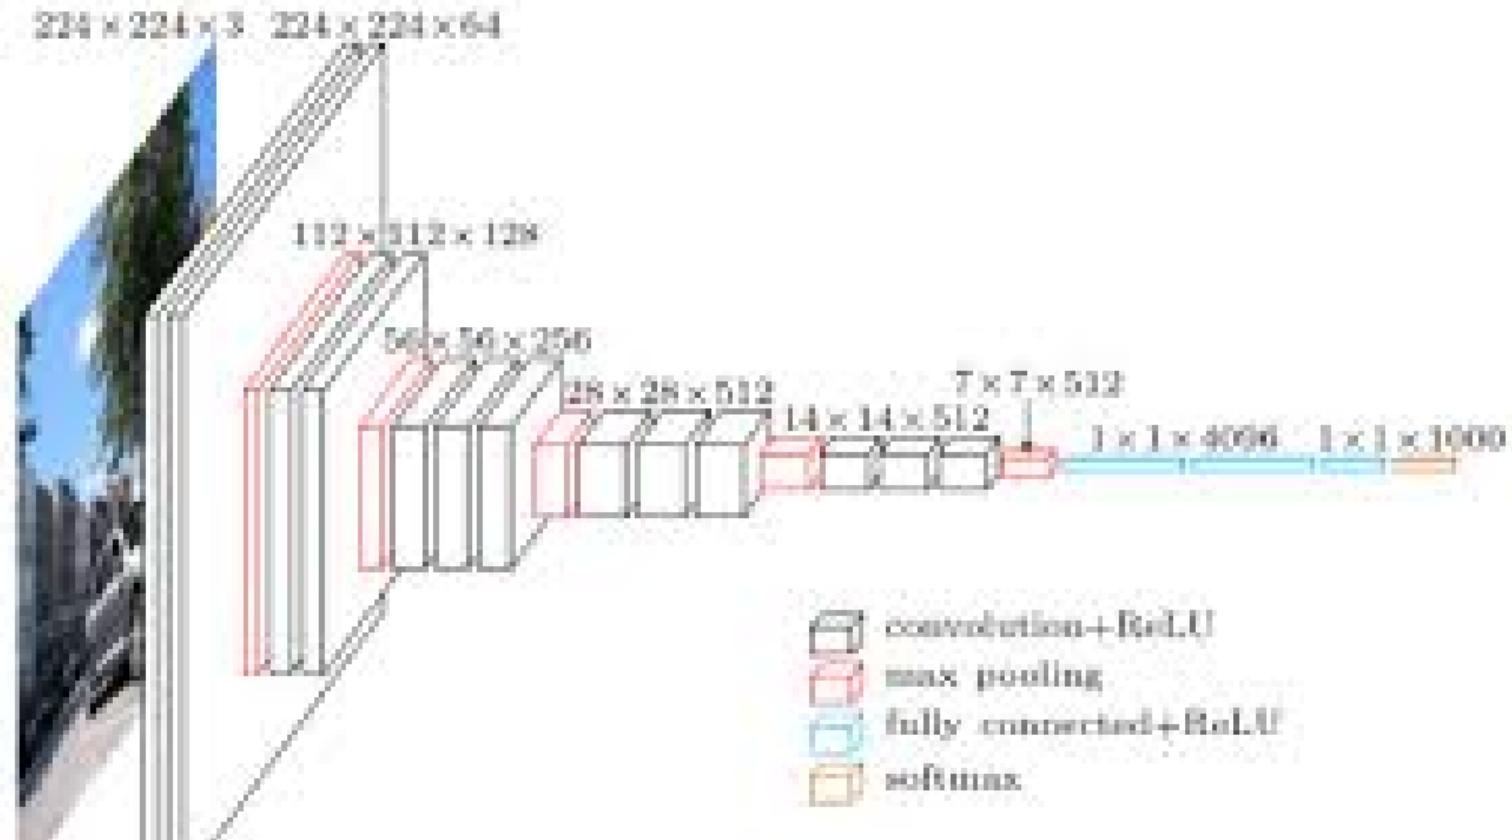


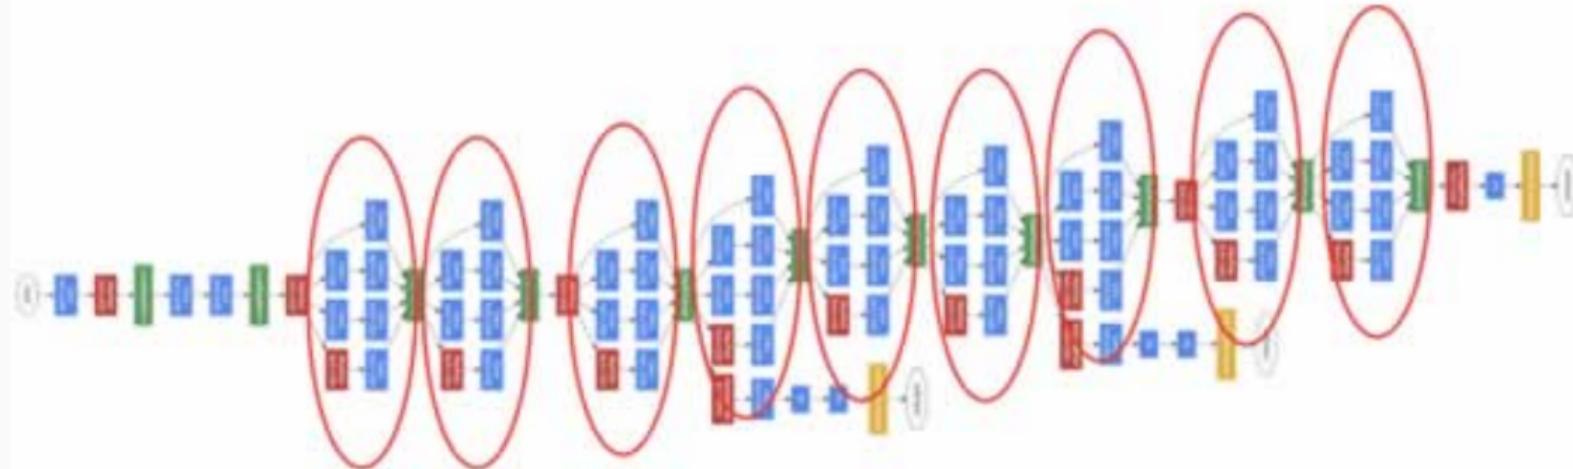
- "The 9 Deep Learning Papers You Need to Know About"  
(<https://adeshpande3.github.io/adeshpande3.github.io/The-9-Deep-Learning-Papers-You-Need-To-Know-About.html>)
- Many of these models are used as benchmarks in deep learning
- There exist many pre-trained versions of the above models that can be leveraged for a new task (this is called **transfer learning**)
- Let's look at a few...

# AlexNet (2012)



# VGG-16 (2014)





## 9 Inception modules

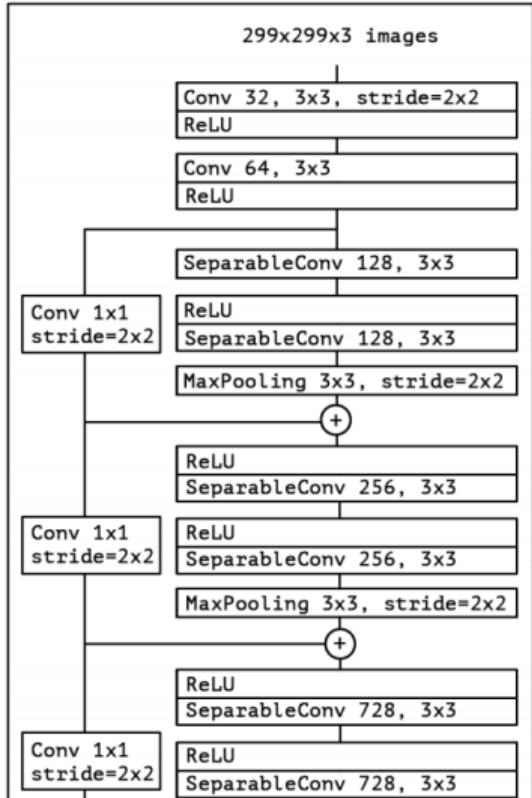
Network in a network in a network...

Convolution  
Pooling  
Softmax  
Other

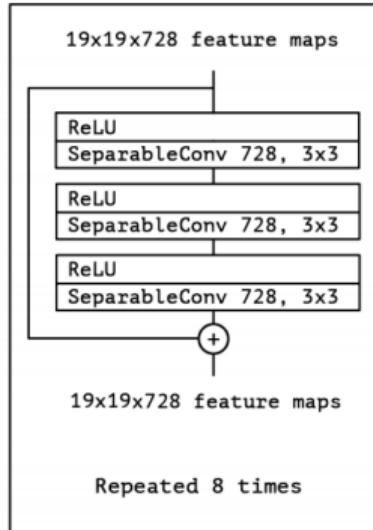
# Xception (2016)



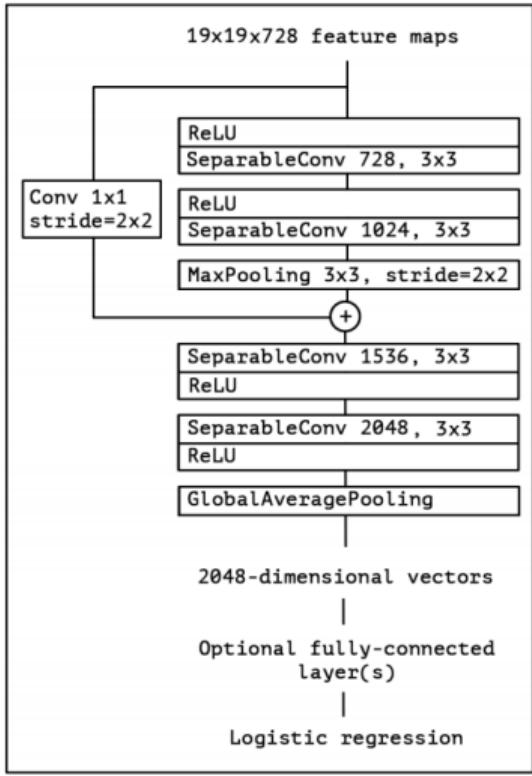
Entry flow



Middle flow



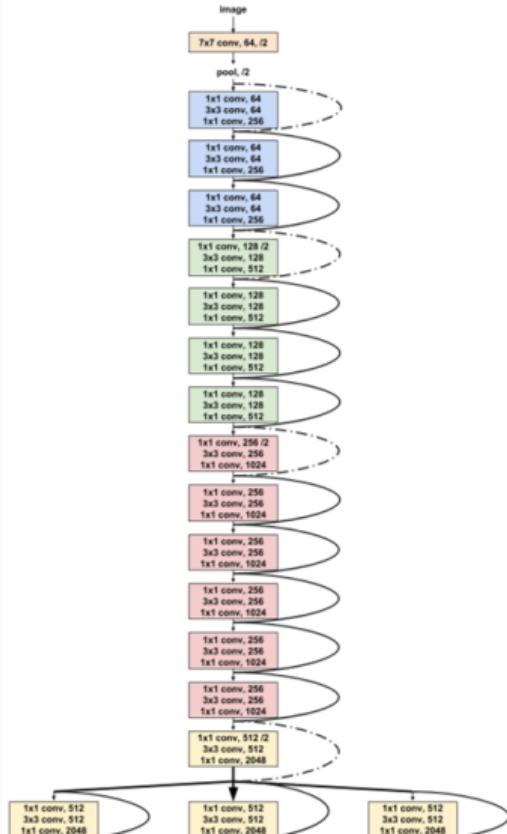
Exit flow



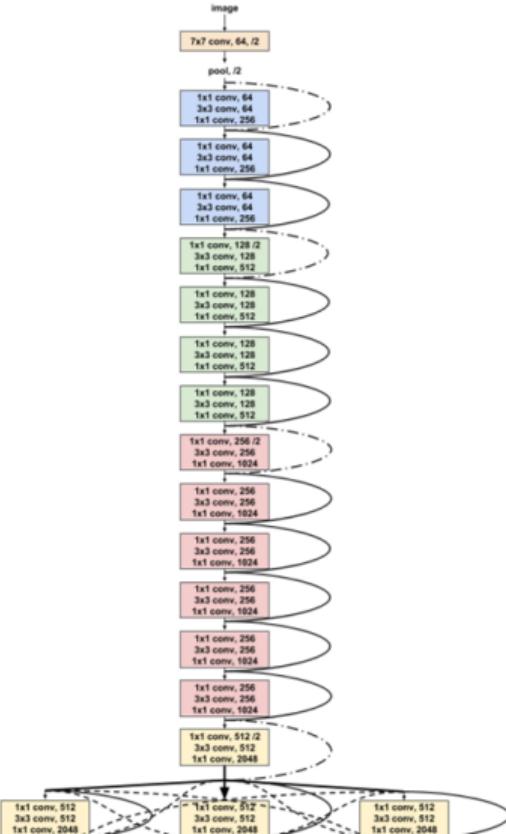
# ResNet-50 (2015)



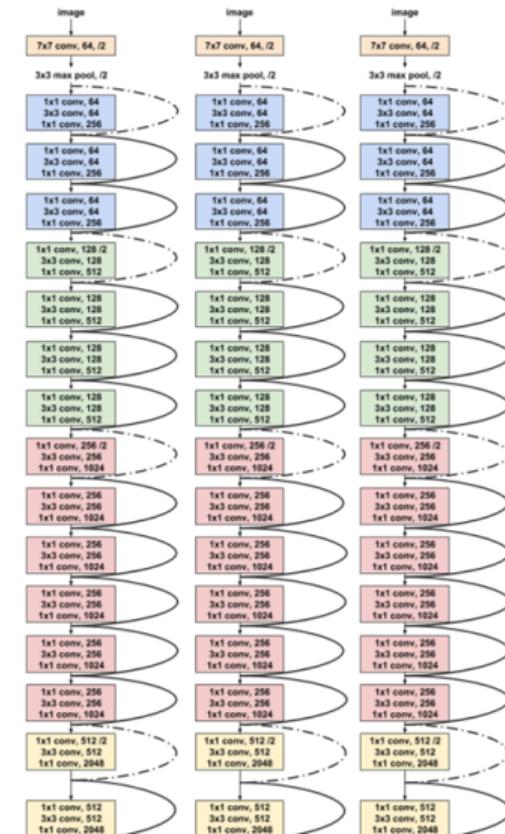
$X_0$ -ResNet-50 (no cross-residuals)



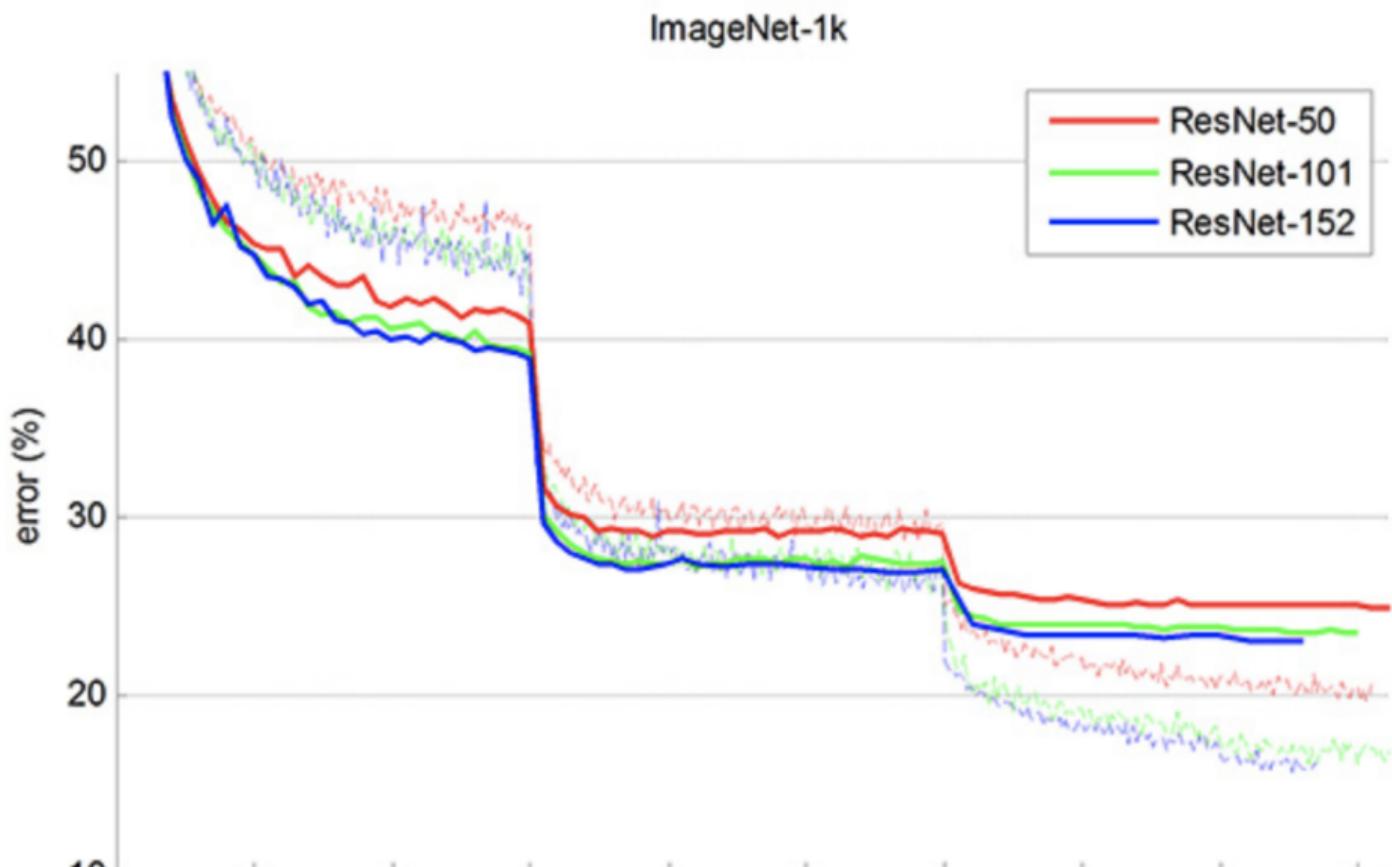
X-ResNet-50 (with cross-residuals)



ResNet-50 (x3)



# ResNet-152 (2015)



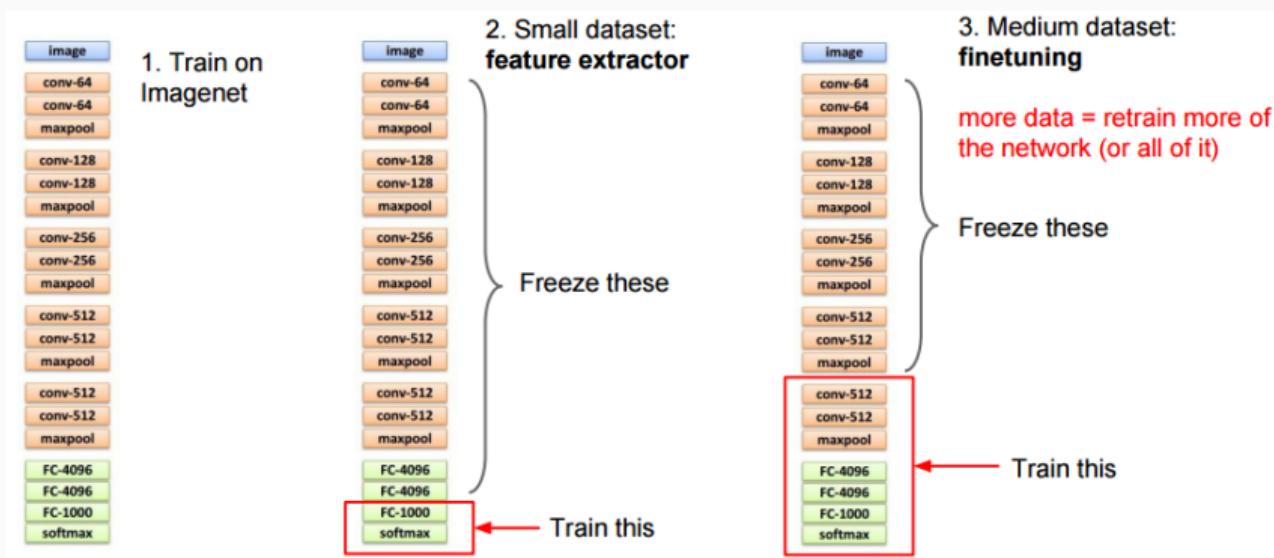
## Transfer Learning



- In practice, we don't train an entire ConvNet from scratch
- One can pretrain the CNN on a large dataset like ImageNet and then finetune the ConvNet so that it does well on our dataset.
- There are several scenarios when we might want to use transfer learning.
- Let's look at a few...



# How Can We Transfer Learn?



# Why does it work?



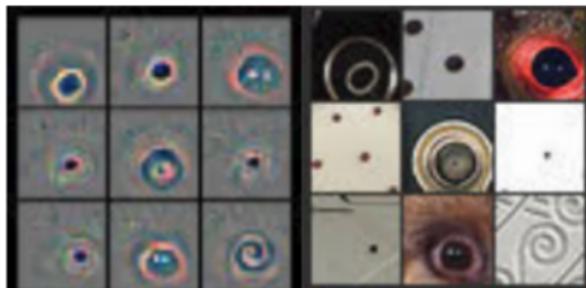
Early layers learn simple features



# Why does it work?



Later layers learn more abstract features



Layer 2 of AlexNet



Layer 4 of AlexNet

## Questions

---

# Questions?

# Demo

---

# CNN demo: CIFAR-10